



Real-time call analytics — architecture (as built)

TypeScript · Node.js 24 + Fastify · React 19 · PostgreSQL 17 · one command: docker compose up · \$0 hosting (Render + Neon)



AI insights — grounded generation pipeline

SQL computes every number · the LLM only selects and phrases · every number in the output is checked against the facts



At 500 enterprise customers — target architecture

~10M calls/day · ~1.5k events/s at peak · ~5k live dashboards · red badges = what breaks first, in order, and where it gets fixed

